

# KFoldCrossValidation

December 20, 2016

## 1 K-Fold Cross Validation

Let's revisit the Iris data set:

```
In [1]: import numpy as np
        from sklearn import cross_validation
        from sklearn import datasets
        from sklearn import svm

        iris = datasets.load_iris()
```

A single train/test split is made easy with the `train_test_split` function in the `cross_validation` library:

```
In [2]: # Split the iris data into train/test data sets with 40% reserved for testing
        X_train, X_test, y_train, y_test = cross_validation.train_test_split(iris.data, iris.target)

        # Build an SVC model for predicting iris classifications using training data
        clf = svm.SVC(kernel='linear', C=1).fit(X_train, y_train)

        # Now measure its performance with the test data
        clf.score(X_test, y_test)
```

```
Out[2]: 0.9666666666666667
```

K-Fold cross validation is just as easy; let's use a K of 5:

```
In [3]: # We give cross_val_score a model, the entire data set and its "real" values, and the number of folds
        scores = cross_validation.cross_val_score(clf, iris.data, iris.target, cv=5)

        # Print the accuracy for each fold:
        print(scores)

        # And the mean accuracy of all 5 folds:
        print(scores.mean())
```

```
[ 0.96666667  1.          0.96666667  0.96666667  1.          ]
0.98
```

Our model is even better than we thought! Can we do better? Let's try a different kernel (poly):

```
In [4]: clf = svm.SVC(kernel='poly', C=1).fit(X_train, y_train)
        scores = cross_validation.cross_val_score(clf, iris.data, iris.target, cv=5)
        print(scores)
        print(scores.mean())

[ 1.          1.          0.9          0.93333333  1.          ]
0.9666666666666667
```

No! The more complex polynomial kernel produced lower accuracy than a simple linear kernel. The polynomial kernel is overfitting. But we couldn't have told that with a single train/test split:

```
In [5]: # Build an SVC model for predicting iris classifications using training data
        clf = svm.SVC(kernel='poly', C=1).fit(X_train, y_train)

        # Now measure its performance with the test data
        clf.score(X_test, y_test)

Out[5]: 0.9666666666666667
```

That's the same score we got with a single train/test split on the linear kernel.

## 1.1 Activity

The "poly" kernel for SVC actually has another attribute for the number of degrees of the polynomial used, which defaults to 3. For example, `svm.SVC(kernel='poly', degree=3, C=1)`

We think the default third-degree polynomial is overfitting, based on the results above. But how about 2? Give that a try and compare it to the linear kernel.

```
In [ ]:
```