

LinearRegression

December 20, 2016

1 Linear Regression

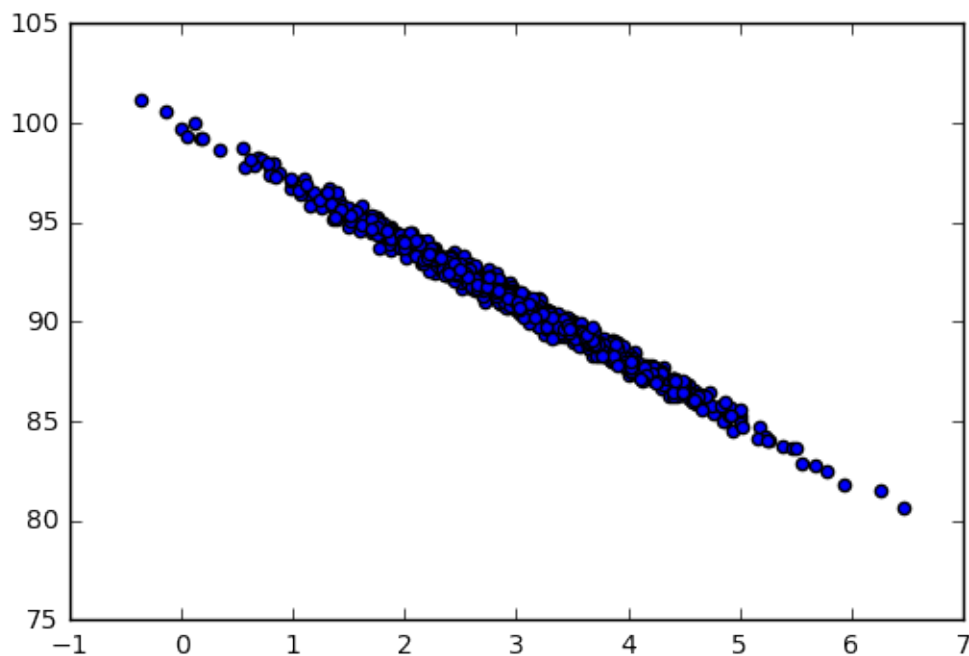
Let's fabricate some data that shows a roughly linear relationship between page speed and amount purchased:

```
In [1]: %matplotlib inline
import numpy as np
from pylab import *

pageSpeeds = np.random.normal(3.0, 1.0, 1000)
purchaseAmount = 100 - (pageSpeeds + np.random.normal(0, 0.1, 1000)) * 3

scatter(pageSpeeds, purchaseAmount)

Out[1]: <matplotlib.collections.PathCollection at 0x5cb8a50>
```



As we only have two features, we can keep it simple and just use `scipy.state.linregress`:

```
In [3]: from scipy import stats
```

```
        slope, intercept, r_value, p_value, std_err = stats.linregress(pageSpeeds, purchaseAmount)
```

Not surprisingly, our R-squared value shows a really good fit:

```
In [4]: r_value ** 2
```

```
Out[4]: 0.98898447900188802
```

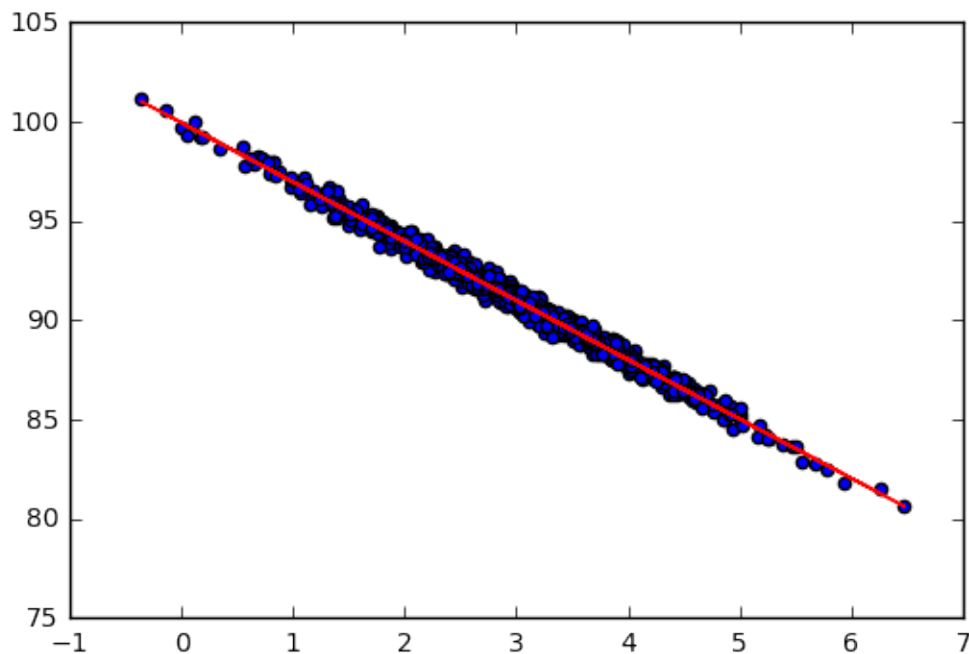
Let's use the slope and intercept we got from the regression to plot predicted values vs. observed:

```
In [5]: import matplotlib.pyplot as plt
```

```
def predict(x):  
    return slope * x + intercept    #  $Y = mX + C$ 
```

```
fitLine = predict(pageSpeeds)    #  $Y$ 
```

```
plt.scatter(pageSpeeds, purchaseAmount)  
plt.plot(pageSpeeds, fitLine, c='r')  
plt.show()
```



1.1 Activity

Try increasing the random variation in the test data, and see what effect it has on the r-squared error value.

In []: